

Java Container Awareness

Version 1.0.0, 2026-08-06

Home

Introduction

Course Title: Java Container Awareness

Description: Course description: this course focus in Java Container details and explanations.

Duration: this course can be finished in 1 to 2 hours

Topics covered in this course

- Introduction to Containerized Java
- Container Resource Detection
- Risks and Memory Management
- OpenShift and Kubernetes Integration
- Troubleshooting and Tooling
- Hands-on Lab

Prerequisites

This course assumes that you have the following prior experience:

- Experience with Java (in general) and heap (off-heap) allocation
- Experience with Openshift in terms of deployment, not necessarily with Operators

Lab Environment

FIXME: Instructions for accessing the lab environment for this hands-on based training.

Introduction to Containerized Java

Containerized Java focuses on the evolution of the Java Virtual Machine (JVM) to operate efficiently within isolated environments like OpenShift and Kubernetes. Modern Java runtimes utilize container awareness to detect cgroup limits, allowing the JVM to automatically adjust its resource consumption based on container constraints rather than host hardware.

Key aspects include:

- **Evolution of Java:** Shifting from host-based resource perception to environment-aware detection.
- **Container Awareness:** The ability of the JVM to recognize and respect memory and CPU limits imposed by the container runtime.
- **Cgroup Detection:** Leveraging kernel-level control groups to monitor and manage resource availability accurately.
- **Operational Efficiency:** Simplifying deployments by ensuring the JVM aligns its internal memory and thread management with defined container resource limits.

This course is fully based on the current OpenJDK implementation and the article [How to use Java container awareness in OpenShift 4](#).

Evolution of Java in containerized environments

The landscape of Java application deployment has shifted dramatically with the advent of containerization. Initially designed for virtual machines and bare-metal hardware, the Java Virtual Machine (JVM) faced significant architectural hurdles when introduced to the world of Linux containers.

Non-container-aware Java

Before 2019, Java (its JVM) would identify the system resources by querying the underlying operating system (OS) kernel via `/proc/meminfo` and `/proc/stat`. When a Java process runs on a physical server, it sees the total memory and CPU cores of the host.

In other words, despite running inside a container, Java was unable to detect the container's limits. It would only detect the host limits regardless of the deployment. Consequently, the user needed to set the heap limits (Xmx) manually as a size.

Therefore, to solve this scenario, it was encapsulated inside an initialization script: `run-java.sh`, which would then detect the container limits and deploy Java with those limits (script → container limits → Java). Java didn't know about cgroups. It would just start, and someone would provide its max heap, Xmx.

JDK8146115

The evolution toward true container awareness reached a milestone with **JDK8146115**. This critical enhancement, backported to Java 8 (starting with 8u191) and native to all newer versions (Java 11,

17, 21+), fundamentally changed how the JVM initializes.

Instead of looking at the host's total resources, the JVM was updated to inspect the cgroup hierarchy directly. This allows the JVM to:

1. **Read Container Memory Limits:** Automatically calculate appropriate initial and maximum heap sizes (`-Xmx` and `-Xms`) based on the specific container memory constraints.
2. **Detect CPU Availability:** Properly identify the number of CPU shares or quotas allocated to the container, allowing the garbage collector and internal thread pools to scale their concurrency appropriately.

Containers achieve isolation through **Control Groups (cgroups)**, which is a Linux kernel feature that limits, accounts for, and isolates the resource usage of a collection of processes. Before 2019, Java was not cgroup-aware when running inside a container, so the JVM would still attempt to query the host's total resources rather than the container's restricted limits. This created a dangerous "mismatch" where the JVM believed it had access to gigabytes of host RAM, leading to:

- **OOMKills:** The JVM would attempt to allocate heap memory based on host capacity, triggering the container orchestrator to kill the process for exceeding its cgroup memory limit.
- **Performance Degradation:** The JVM would size its internal thread pools (e.g., ForkJoinPool) based on the total physical CPU cores of the host, leading to excessive context switching and performance bottlenecks in environments where the container was throttled, given the container likely has different CPU limits.

Current Fully Container-Aware Java

Today, container awareness is the standard expectation for cloud-native Java development for Java 8, 11, 17, 21, and 25. The transition from unaware to fully container-aware has considerably streamlined the DevOps workflow:

- **Self-Tuning:** Applications can now adjust their memory footprint dynamically to fit within Kubernetes/OpenShift resource requests and limits without manual intervention.
- **Improved Stability:** By aligning the JVM's view of resources with the kernel-enforced limits, the frequency of container restarts due to resource exhaustion is drastically reduced.
- **Reduced Configuration Overhead:** Developers do not spend time tuning CPU values and `-Xmx` limits, as the JVM can now effectively derive them from the container environment.

Understanding this evolution is crucial for any team moving toward microservices architectures. Container awareness provides a robust foundation that requires the development team (who handles the Java settings) and the Deployment team (who handles the container size) to balance JVM heap memory, off-heap memory, and the container's hard limits—a topic that remains a critical consideration for production-grade deployments.

Overview of container-aware versus non-container-aware Java

In the early days of containerization, Java applications often struggled to coexist with container runtimes like Docker and Kubernetes. This friction stemmed from a fundamental mismatch: the Java Virtual Machine (JVM) was designed to query the host operating system for resource

availability, while containers rely on control groups (cgroups) to impose artificial resource constraints.

The Evolution of Container Awareness

Before the introduction of container awareness, the JVM would query the host machine's total physical memory and CPU cores, completely ignoring the limits imposed by the container runtime.

Feature	Non-Container-Aware (Legacy)
Fully Container-Aware (Current)	Memory Detection
Sees total host RAM	Sees cgroup-defined memory limit
CPU Detection	Sees total host CPU cores
Sees cgroup-defined CPU quota	Heap Sizing
Based on total host memory	Based on container memory limit
OOM Risk	High (Host-level OOM)

Non-Container-Aware Java: The Visibility Gap

In a non-container-aware environment, if you deploy a Java application with a default heap size (e.g., `-Xmx` not explicitly defined) into a container with a 1GB limit, the JVM might inspect the host's 64GB of RAM and attempt to allocate a heap significantly larger than 1GB.

When the application tries to consume this memory, the container runtime's kernel-level cgroup enforcement triggers an **Out-of-Memory (OOM) Kill**. The process is terminated abruptly because it breached the container's allowed boundary, despite the JVM "thinking" it had plenty of memory available.

Container-Aware Java: The Shift with JDK8146115

The industry shift occurred with the implementation of **JDK8146115**, starting with Java 8u191 and continuing through all modern LTS releases (e.g., Java 11, 17, 21). This update enabled the JVM to "read" the `cgroup` filesystem directly.

When running a container-aware JVM, the process performs the following:

- * **Memory Sensing:** Instead of asking the OS for total physical RAM, the JVM reads the `memory.limit_in_bytes` file from the container's cgroup. It then calculates the default heap size as a percentage of this limit, rather than a percentage of the host's total RAM.
- * **CPU Awareness:** The JVM queries the `cpu.shares` or `cpu.cfs_quota_us` files to determine how many CPU cycles are available to the container. This prevents the JVM from spawning an excessive number of Garbage Collection (GC) threads that would otherwise lead to performance degradation via context switching.

Why This Matters for Modern Deployments

Transitioning to container-aware Java is not merely a technical optimization; it is a prerequisite for stability in orchestrators like OpenShift and Kubernetes.

- **Simplified Configuration:** Developers no longer need to hardcode `-Xmx` values that perfectly

match every environment's resource limit. The JVM automatically adapts to the **limits** defined in your Kubernetes deployment manifests.

- **Deterministic Behavior:** By aligning the JVM's view of resources with the container's reality, you reduce the likelihood of "invisible" OOM kills that occur when the JVM over-allocates memory.
- **Operational Efficiency:** Container-aware Java allows DevOps teams to set resource requests and limits in OpenShift with the confidence that the application will scale its internal memory management to fit the provided space.



While container awareness resolves around JVM limits association with the container size, such as heap and threads, it does not necessarily account for **off-heap** memory (e.g., metaspace, code cache, direct buffers, or native memory). Therefore, even with container aware Java, one must benchmark your application to ensure that the ratio of heap to off-heap memory is sufficient resources are allocated to avoid container-level OOMkills events.

Importance of cgroup detection for deployments

In cloud-native architectures, the shift from virtual machines to containerized runtimes has fundamentally changed how applications interact with hardware. For Java, historically accustomed to seeing the "full capacity" of a host OS, the introduction of cgroups (control groups) represents a critical evolution in resource management.

The Evolution of Java and Resource Visibility

Before container-aware JVMs, a Java application running inside a container would often query the underlying host OS to determine available memory and CPU. If the container was restricted to 2GB of memory, but the host node had 64GB, the JVM would frequently miscalculate its maximum heap size (**-Xmx**). This led to two primary failure modes:

1. **OOMKiller Invocation:** The JVM, believing it had access to a larger pool of memory, would allocate heap space that exceeded the container's hard limit. The Linux kernel, seeing the container exceed its cgroup limit, would trigger an Out-of-Memory (OOM) Kill, abruptly terminating the process.
2. **Resource Inefficiency:** Conversely, if the JVM underestimated the capacity, it might fail to utilize the resources explicitly allocated to it by the orchestrator, leading to poor performance and degraded throughput.

Why cgroup Detection is Critical

Cgroup detection is the mechanism by which the JVM "understands" the boundaries imposed by the container runtime. Its importance for deployments cannot be overstated for the following reasons:

- **Boundary Compliance:** It ensures the JVM respects the container's quota rather than the physical host's capacity. By reading from **/sys/fs/cgroup/** (or equivalent cgroup v2 paths), the JVM adjusts its internal ergonomics to fit within the **cgroup** limits.

- **Preventing "Silent" Failures:** With proper cgroup detection, the JVM aligns its garbage collection triggers and heap allocation thresholds with the container's memory limit. This prevents the scenario where the application performs perfectly under light load but crashes instantly when the memory pressure hits the cgroup ceiling.
- **Orchestration Alignment:** Orchestrators like OpenShift and Kubernetes use cgroups to enforce Resource Limits and Requests. Java's ability to "see" these limits allows it to cooperate with the scheduler. When the JVM is container-aware, it can effectively use flags like `-XX:MaxRAMPercentage`, ensuring the heap scales dynamically with the pod's assigned limit.

Implications for DevOps teams

For a DevOps team, relying on container-aware Java simplifies the deployment pipeline. You no longer need to hardcode `-Xmx` values for every environment. By allowing the JVM to detect the cgroup limit, you can:

- **Standardize Images:** Use the same container image across development, testing, and production environments, where resource limits may vary significantly.
- **Predictable Scaling:** As you adjust memory limits in your Kubernetes manifest, the application automatically adapts its memory footprint without requiring a rebuild or manual configuration update.



While cgroup detection prevents the JVM from over-allocating heap space, it does **not** manage off-heap memory (Metaspace, Code Cache, Thread stacks, etc.). Because the heap is only a portion of the total memory usage, administrators must still account for the "non-heap" overhead. Failing to leave enough room for off-heap memory can still result in a cgroup OOMKill, even if the heap appears to be within the allowed percentage.

By ensuring your deployments utilize container-aware Java versions (such as Java 8u191+ or any later LTS version), you bridge the gap between the application's runtime needs and the platform's enforcement mechanisms, leading to more resilient and stable microservices.

Container Resource Detection

Container resource detection enables the Java Virtual Machine (JVM) to accurately identify resource constraints imposed by the container environment rather than the underlying host node.

- **JVM Evolution:** Starting with Java 8u191 (via JDK8146115), the JVM gained the ability to parse control group (cgroup) files to determine resource boundaries.
- **Memory Limit Detection:** The JVM identifies the specific memory limit assigned to the container, allowing it to scale the heap size appropriately based on available container resources.
- **CPU Availability:** The JVM detects the number of CPU cores allocated to the container, preventing over-provisioning and ensuring efficient thread management within the containerized environment.
- **Operational Benefits:** Accurate detection prevents "out-of-memory" errors caused by the JVM incorrectly defaulting to host-level physical memory rather than container-specific limits.

Understanding JDK8146115 and Java 8u191 Updates

Before the introduction of specific container-aware updates, the Java Virtual Machine (JVM) lacked the ability to natively "see" the resource boundaries imposed by container runtimes. This caused significant operational friction in cloud-native environments.

The Challenge: JVM Resource Blindness

In early containerized environments, the JVM would query the underlying host operating system for available resources (such as total system memory and available CPU cores) rather than the specific cgroup limits applied to the container.

This behavior led to a critical disconnect:

- * **Memory Miscalculation:** The JVM would attempt to size its heap based on the host's physical RAM, often exceeding the container's hard memory limit set by Kubernetes or OpenShift.
- * **OOM Killer Intervention:** Once the Java process exceeded the container's cgroup memory limit, the Linux kernel's Out-Of-Memory (OOM) Killer would terminate the process, leading to unexpected pod restarts.
- * **CPU Throttling:** Similar issues occurred with CPU calculation, where the JVM would perceive all host cores as available, leading to inefficient thread pool sizing and potential context-switching overhead.

The Turning Point: JDK8146115

To address these issues, the OpenJDK community introduced **JDK8146115**. This enhancement provided the necessary infrastructure for the JVM to correctly read and respect cgroup memory and CPU limits.

Instead of querying `/proc/meminfo` or `/proc/stat`—which reflect host-level resources—the JVM was updated to inspect the cgroup filesystem (`/sys/fs/cgroup/`) directly. This allows the JVM to identify the specific constraints applied to the process namespace.

Java 8u191: The Paradigm Shift

While JDK8146115 provided the foundational logic, the release of **Java 8u191** served as the catalyst for widespread adoption. This update backported the container-awareness features, effectively bringing modern cloud-native capabilities to long-standing Java 8 applications.

Key changes introduced in 8u191 include:

- * **Native Cgroup Awareness:** The JVM now automatically detects container limits at startup.
- * **Default Ergonomics:** The default heap size calculation logic was updated to use a fraction of the container's memory limit instead of the host's physical memory.
- * **New JVM Flags:** Introduction of flags such as `-XX:MaxRAMPercentage`, `-XX:MinRAMPercentage`, and `-XX:InitialRAMPercentage`, allowing developers to define heap sizes as a ratio of the container limit rather than hard-coded values.

Impact on Modern Deployments

With Java 8u191+, the relationship between the container and the JVM became **coupled**.

Feature	Legacy Behavior (Pre-8u191)
Modern Behavior (8u191+)	Memory Source
Queries <code>/proc/meminfo</code> (Host RAM)	Queries cgroups (Container Limit)
Heap Sizing	Requires manual <code>-Xmx</code> configuration
Automatically scales based on <code>-XX:MaxRAMPercentage</code>	Operational Impact
Decoupled; changing container size requires manual JVM tuning	Coupled; scaling container limits automatically adjusts heap

By aligning the JVM's view of resources with the container orchestration layer, these updates eliminated the "decoupling" problem. Developers no longer need to manually synchronize `-Xmx` settings with pod resource limits, significantly simplifying the management of Java applications in Red Hat OpenShift and Kubernetes.

Detecting Container Memory Limits

In modern Java deployments, it is critical that the Java Virtual Machine (JVM) correctly identifies the memory boundaries imposed by the container orchestration platform. Failure to do so leads to performance degradation, inefficient resource utilization, or unexpected pod terminations.

Why Container Awareness Matters

Prior to the integration of **JDK8146115** (introduced in Java 8u191), the JVM relied on standard OS-level utilities to detect available memory. In a containerized environment, these utilities often reported the total memory of the underlying physical host (the OCP node) rather than the restricted memory assigned to the specific container.

If the JVM perceives the node's total memory instead of the container's limit, it may attempt to allocate a heap size that exceeds the container's allowed threshold, resulting in an `OutOfMemoryException` or, more commonly, a `SIGKILL` by the Linux kernel's Out-Of-Memory (OOM)

killer.

How the JVM Detects Container Limits

When running in a container-aware environment, the JVM inspects the control group (cgroup) filesystem. By reading the files exposed within the container—specifically those under `/sys/fs/cgroup/memory`—the JVM determines the hard limits imposed by the container orchestrator.

Verification Tools

To determine if your Java application is correctly "seeing" the container limits, you should utilize built-in JVM diagnostic flags rather than relying on standard Linux utilities like `free -m`.



Standard utilities like `free`, `top`, or `lscpu` are **not container-aware**. They will report the resources of the host node, leading to incorrect diagnostic conclusions. Use container-aware tools, such as `jcmd VM.info`, and `-XshowSettings:system`.

Using `-XshowSettings:system`

The most direct way to verify what the JVM has detected is to pass the `-XshowSettings:system` flag during application startup or via an environment variable.

Example Output:

```
Operating System Metrics:
  Provider: cgroupv1
  Memory Limit: 128.00M
  Memory Soft Limit: unlimited
  Memory Current: 45.21M
  CPU Shares: 512
  CPU Quota: 100000
  CPU Period: 100000
```

When the JVM is container-aware, you will see explicit lines reflecting the **Memory Limit**. If this value matches your OpenShift/Kubernetes request/limit configuration, the JVM is successfully performing container detection.

Troubleshooting Memory Misalignment

If your application exhibits the following symptoms, the JVM is likely failing to respect container limits:

1. **Host-level reporting:** Non-container-aware commands executed inside the container (e.g., `free -m`) show memory values matching the physical node rather than the pod limit.
2. **OOM Killer termination:** The pod restarts with an exit code indicating OOM, despite the `-Xmx` setting appearing reasonable for the application's perceived workload.

Two main scenarios:

1. **Underutilization:** The application heap is significantly smaller than the container limit because it defaulted to a generic fraction of the host's massive memory footprint, leading to frequent Garbage Collection (GC) pauses.
2. **Overutilization:** The application heap (plus off-heap) is significantly larger than the container limit, which will eventually lead to OOMKill events.

Diagnostic Strategy

To analyze resource constraints effectively:

- * **Check the logs:** Inspect your pod logs to confirm if `run-java.sh` or the standard OpenJDK startup scripts are being utilized, as these often contain logic to calculate heap sizes based on container limits.
- * **Verify jcmd output:** Use `jcmd <pid> VM.info` to inspect the runtime configuration and ensure the JVM has correctly parsed the `cgroup` memory controller.
- * **Compare with OCP Metadata:** Always compare the JVM's reported `Memory Limit` against the `resources.limits.memory` configured in your OpenShift `DeploymentConfig` or `Deployment` YAML.

Determining CPU Availability Within Containers

In a virtualized or bare-metal environment, the Java Virtual Machine (JVM) traditionally queries the operating system for the total number of physical CPUs available on the host. However, in containerized environments like Red Hat OpenShift or Kubernetes, a container is typically restricted to a subset of the host's compute resources.

If the JVM is not container-aware, it may incorrectly identify the host's total CPU count rather than the container's allocated limit. This mismatch can lead to inefficient thread pool sizing (such as `ForkJoinPool` or parallel garbage collection threads), causing contention and performance degradation.

How Java Detects CPU Resources

Modern JDKs (Java 8u191+ and all subsequent versions) use **cgroups (Control Groups)** to determine CPU availability. The JVM reads the cgroup filesystem to calculate the `ActiveProcessorCount`.

The Calculation Logic

When running inside a container, the JVM inspects two primary metrics from the cgroup controller to determine how many processors it should use:

1. **CPU Shares (`cpu.shares`):** This represents the relative weight of the container. If the host has 64 cores and your container has a share weight that equates to 2 cores, the JVM uses this to influence its internal parallelism.
2. **CPU Quota and Period (`cpu.cfs_quota_us` and `cpu.cfs_period_us`):** These define a hard limit. For example, if the quota is 200,000 and the period is 100,000, the container is limited to 2 cores ($\$200,000 / 100,000 = 2\$$).

The JVM calculates the number of processors as the minimum of the host's processors and the calculated container limit.

Diagnostic Tools: Container-Aware vs. Non-Aware

To verify how your Java application perceives CPU availability, you must distinguish between tools that see the "Host" view versus the "Container" view.

Tool Category	Examples
Insight	Non-Container Aware
<code>lscpu, nproc</code>	Returns the total CPU count of the underlying OpenShift Node . This is often misleading for capacity planning.
Container Aware	<code>VM.info, -XX:+PrintFlagsFinal</code>

Verifying CPU Detection

You can inspect the JVM's view of your container's resources by using the `jcmd` utility.

1. Access your pod terminal: `oc rsh <pod-name>`
2. Run `jcmd` to print VM information: `jcmd <pid> VM.info`

Look for the `container` section in the output. A container-aware JVM will report output similar to this:

```
container (cgroup) information:
container_type: cgroupv1
cpu_cpuset_cpus: 0-127
cpu_memory_nodes: 0-7
active_processor_count: 1  <-- JVM correctly sees 1 CPU limit
cpu_quota: -1
cpu_period: 100000
cpu_shares: 51
```

Best Practices for CPU Configuration

- **Define Requests and Limits:** Always explicitly define `resources.limits.cpu` in your OpenShift deployment YAML. This ensures the cgroup values are populated correctly.
- **Avoid Oversizing:** If you set a high CPU limit but your application is single-threaded, the JVM may spawn excessive GC threads, leading to context-switching overhead.
- **Monitor `active_processor_count`:** If `active_processor_count` does not match your expectations, ensure you are running a supported OpenJDK version (8u191 or higher). If you are forced to use an older, non-aware version, you must manually set `-XX:ActiveProcessorCount` to match your container limit.

Hands-on Exercise: Inspecting CPU Awareness

In this short task, you will verify the JVM's CPU detection.

1. **Inspect current CPU perception:** Inside your running container, identify the Java process ID:
`jps`
2. **Query the VM metadata:** Run the following command to see the active processor count: `jcmd <pid> VM.info | grep active_processor_count`
3. **Compare with OS-level tools:** Run `nproc` in the same terminal. **Note: If `nproc` returns the node's total CPUs (e.g., 64) but `active_processor_count` returns your container limit (e.g., 2), your application is correctly container-aware.**

Risks and Memory Management

When deploying Java in containerized environments, improper memory configuration can lead to significant stability and performance issues. Understanding the balance between heap and off-heap memory is critical to preventing infrastructure-level failures.

- **Decoupling Risks:** Hard-coding heap sizes (e.g., via `-Xmx`) without considering container limits can lead to memory mismatches. If the JVM ignores container boundaries, it may trigger cgroups OOM-Kills, causing the kernel to terminate the application.
- **Off-Heap Considerations:** Container-aware Java calculates heap size as a percentage of available container memory. It is essential to account for off-heap usage (e.g., metaspace, thread stacks, native buffers) to ensure the total memory footprint does not exceed the container limit.
- **Resource Underutilization:** Inadequate configuration can lead to inefficient resource usage, where the JVM reserves significantly less memory than the container provides, leaving allocated infrastructure resources idle.
- **Best Practices:**
 - Prioritize container-aware JVM settings over manual heap overrides to allow for dynamic scaling.
 - Benchmark application memory usage to determine the optimal ratio between heap and off-heap requirements.
 - Avoid configurations that ignore cgroup limits, as these bypass the self-tuning capabilities of modern JDKs.

Implications of Decoupling Heap Memory from Container Limits

Inside a containerized environment, the JVM's ability to detect container resource boundaries is a critical feature. When developers manually specify heap sizes using the traditional heap limit: `-Xmx` without regard for the container's cgroup limits, they risk decoupling the JVM memory footprint from the orchestration layer. This decoupling often leads to significant operational instability and an inherited risk of OOMKill.

The Risks of Manual Overrides

Before Java became fully container-aware, it was common practice to explicitly define `-Xmx` (maximum heap size) to ensure the JVM did not attempt to consume all available memory on a physical host. For example, on a straight VM deployment. However, in a containerized world, this approach creates several hazards:

- **cgroup OOMKills:** If the `-Xmx` is set too high relative to the container's memory limit, the total memory usage—heap plus off-heap (metaspace, thread stacks, GC structures, and native buffers)—will eventually exceed the cgroup limit. The container runtime or the kernel will detect this violation and trigger an `OutOfMemoryKill` (OOMKill), terminating the process abruptly.
- **Resource Underutilization:** If the `-Xmx` is set significantly lower than the container limit, the

application cannot scale its memory usage to handle increased load, even if the container has plenty of headroom. This results in inefficient resource allocation and potential `OutOfMemoryException` (OOM) errors within the JVM, despite the container appearing to have available memory.

- **Off-heap limits:** DevOps teams/developers may often focus exclusively on the heap. However, native memory (off-heap) usage is not governed by `Xmx`, although `Xmx` caps the heap limit. The JVM is a combination of both heap and off-heap, and both segments together potentially cause the combined memory pressure to breach the container's hard limit.

Why Decoupling is an Anti-Pattern

Decoupling occurs when the JVM's memory configuration ignores the orchestrator's resource constraints. This makes the application "blind" to the environment it resides in.

Modern Java versions (Java 8u191+, Java 11+) are designed to be container-aware, meaning they automatically derive memory settings from cgroups. Overriding this by hardcoding heap values removes the benefits of this feature, essentially reverting the JVM to a "legacy" mode that is unaware of its surroundings.

Best Practices: The Percentage-Based Approach

To avoid decoupling, the recommendation is to shift from absolute values (`-Xmx`) to relative values. By using flags like `-XX:MaxRAMPercentage`, the JVM maintains a "coupled" relationship with the container:

- **Dynamic Changes:** As you increase the memory limit of your OpenShift or Kubernetes pod, the JVM heap will not automatically adjust to maintain the defined percentage. That's because the JVM only detects the limits once when it starts and therefore needs to fully restart to detect its limits.
- **Safety Margin:** Leaving a percentage of the container memory free (i.e., not allocating 100% to the heap) provides the necessary space for off-heap components—such as Netty buffers, thread stacks, and Metaspace—preventing cgroup-level OOMKills.



If you must explicitly define memory, use `-XX:MaxRAMPercentage=N`. This ensures the heap size is always a derivative of the container limit, preserving the coupling between the JVM and the infrastructure.

Summary of Implications

Risk Factor	Consequence	Impact
Fixed <code>-Xmx</code>	Decoupling from cgroups	Unexpected OOMKills
Ignored Off-Heap	Native memory starvation	JVM instability / Crashes
Manual Constraints	Operational rigidity	Need for manual tuning during scaling

By allowing the JVM to detect container limits automatically, i.e. by utilizing percentage-based

configurations, the DevOps teams ensure the Java application remains dynamic, resilient, portable, and correctly scaled within your cluster limits and boundaries.

Managing Off-Heap Memory in Containerized Environments

In modern Java applications, memory management extends beyond the Java Heap. While the Heap is responsible for storing object instances, the **Off-Heap** memory—encompassing Metaspace, Code Cache, Thread stacks, and Direct ByteBuffers—is critical for the stability of containerized deployments.

When a Java application runs in an OpenShift or Kubernetes container, it is subject to strict Linux Control Group (cgroup) memory limits. If the sum of the Heap and Off-Heap memory exceeds these container limits, the container runtime will trigger an **OOMKill** (Out of Memory Kill), regardless of how much space is left in the Java Heap.

The Decoupling Risk

Historically, developers focused exclusively on **-Xmx** (maximum heap size). However, in a container, the JVM sees the container's memory limit. If you configure the Heap to consume too large a percentage of the container's limit, you leave insufficient headroom for Off-Heap requirements.

Key Off-Heap components that must be accounted for include:

- **Metaspace:** Stores class metadata.
- **Code Cache:** Stores compiled native code by the JIT compiler.
- **Thread Stacks:** Each thread requires native memory; in applications with high thread counts, this can lead to significant memory consumption.
- **Direct Buffers:** Used by NIO libraries and netty-based frameworks.

Best Practices for Configuration

To prevent **OOMKills**, you must treat the container's memory limit as the total capacity for both Heap and Off-Heap.

1. **Calculate the Total Footprint:** Before setting heap percentages, perform load testing to determine the "baseline" Off-Heap usage of your application under peak load.
2. **Use Relative Sizing:** Instead of fixed **-Xmx** values, leverage container-aware JVM flags that allow the heap to scale proportionally to the container limit.

```
-XX:MaxRAMPercentage=70.0
```

Note: If your application is highly dependent on Off-Heap (e.g., uses large Direct Buffers or high thread counts), you should reduce this percentage to 50% or 60% to ensure enough "breathing room" for native memory allocations.

3. **Monitor Native Memory:** Use the Native Memory Tracking (NMT) feature to gain visibility into non-heap memory usage.

```
# Enable NMT in the container
-XX:NativeMemoryTracking=summary
```

Hands-on Activity: Analyzing Off-Heap Requirements

This activity guides you through identifying the native memory footprint of your Java process.

Step 1: Enable NMT and Inspect Memory

Use `jcmd` to inspect the memory distribution of your running containerized application.

1. Access the terminal of your running pod:

```
oc rsh <pod-name>
```

2. Run the NMT summary command:

```
jcmd <pid> VM.native_memory summary
```

3. Analyze the output. Look specifically at:
 - **Internal:** Memory used by the JVM internal structures.
 - **Symbol:** Memory used for string interning.
 - **Thread:** Memory consumed by stack spaces.

Step 2: Simulate Constraint Pressure

To ensure your configuration is resilient, verify that your application does not exceed the container limit.

1. View the current cgroup memory limit for your container:

```
cat /sys/fs/cgroup/memory/memory.limit_in_bytes
```

2. Compare this value against your calculated Heap + Off-Heap usage. If your application's **Total** committed memory (as seen in NMT) + `-Xmx` exceeds this limit, you must adjust your `MaxRAMPercentage` downward or increase the container's memory request/limit in your OpenShift `DeploymentConfig`.

Troubleshooting Tip

If you encounter `cgroup OOMKills` despite having a small Java Heap, use `jcmd` to identify if the

Thread or **Internal** categories are consuming unexpected amounts of native memory. If so, investigate thread pool sizing or consider decreasing the `-Xss` (thread stack size) flag to reduce the memory footprint per thread.

Best Practices for Java Memory Configuration in Containers

As Java has evolved to be flexible and container-aware, some practices could be considered anti-patterns, such as:

- `Xmx` direct usage: in most cases, it is an anti-pattern. In other words, the manual strategy of defining absolute heap sizes (e.g., `-Xmx2g`) has become a deployment anti-pattern, because it decouples from the container size.
- JVM arguments directly at build time, rather than runtime. For example, the usage of JVM arguments hardcoded in a Containerfile (Dockerfile), rather than customizable in a deployment resource (YAML).

When the JVM is unaware of its container boundaries, it will default to host-level metrics, leading to resource contention, poor performance, or premature `cgroups` OOM-Kills.

The modern approach focuses on **relative size configuration** to maintain a healthy balance between the JVM heap and the native (off-heap) memory required by the container environment, so the size of Java's heap changes based on the container size (dynamically).

Moving from Absolute to Relative Configuration

To ensure your Java application scales gracefully within an OpenShift or Kubernetes pod, follow these best practices for memory management.

1. Use Percentage-Based Flags

Instead of hardcoding memory values, utilize percentage-based flags. These ensure the JVM dynamically adjusts its heap size relative to the memory limits defined in your Pod specification.

- `-XX:MaxRAMPercentage=N`: Defines the maximum heap size as a percentage of the container's memory limit. For most applications, a value between `50.0` and `75.0` is recommended, leaving the remaining space for off-heap requirements.
- `-XX:InitialRAMPercentage=N`: Sets the initial heap size as a percentage of the container limit.
- `-XX:MinRAMPercentage=N`: Sets the minimum heap size when the container limit is very small.

2. Understanding the Off-Heap "Gap"

The memory not occupied by the JVM heap is not "wasted"—it is critical for the stability of the container. When you configure your heap, you must leave sufficient headroom for:

- **Metaspace**: Stores class metadata.
- **Code Cache**: Stores JIT-compiled native code.

- **Thread Stacks:** Each thread requires native memory; a large number of threads can lead to unexpected native memory pressure.
- **Direct Byte Buffers:** Often used by frameworks like Netty, AMQ, or Data Grid.
- **Garbage Collector Metadata:** The structures required by G1GC or ZGC to track object references.

Best Practice: If your application relies heavily on off-heap memory (e.g., streaming or high-throughput messaging), use a lower `MaxRAMPercentage` (e.g., 50%) to prevent the total container usage from hitting the `cgroups` limit.

3. Avoid MaxRAM Override

Avoid setting the `-XX:MaxRAM` flag. When you set `MaxRAM` manually, you effectively decouple the JVM's view of memory from the actual container `cgroups` limit. This defeats the purpose of container awareness and can lead to the JVM attempting to allocate more memory than the container is permitted to use, resulting in an immediate OOM-Kill by the kernel.

4. Aligning with Container Orchestration

Ensure your Pod resource requests and limits are configured to support your Java workload:

- **Memory Limits:** Always set a memory limit in your Deployment/DeploymentConfig. The JVM uses this value to calculate the heap size when using percentage-based flags.
- **Avoid "Oversizing":** Do not set the limit significantly higher than the application's actual footprint. While Kubernetes handles overcommitment, the JVM will calculate its heap based on the **limit**, not the actual usage.

Summary Table: Configuration Comparison

Approach	Recommendation
<code>-Xmx</code> (Absolute)	Discouraged. Causes decoupling and potential OOM-Kills if container limits change.
<code>-XX:MaxRAMPercentage</code>	Recommended. Automatically aligns heap with container <code>cgroups</code> limits.
<code>MaxRAM</code>	Avoid. Can create conflicting memory views between the OS and the JVM.



When using Red Hat OpenJDK images, the provided `run-java.sh` initialization script sets a specific `MaxRAMPercentage` percentage, currently at 80%. Previously, the `MaxRAMPercentage` was 50%. This value sets the Max Heap at 80% of the container size. The remainder goes to off-heap. Overriding the `MaxRAMPercentage` with a custom value is sufficient to customize the deployment, considering that 80% heap may not be adequate.

Troubleshooting Resource Constraints

If your application is experiencing `OutOfMemoryException` or being killed by the container orchestrator, analyze the environment with the following diagnostic flow:

1. **Check VM.info file:** Use Java Diagnostics (jcmd) to print all details of Java concerning containers.
2. **Check cgroups usage:** Use `cat /sys/fs/cgroup/memory/memory.stat` inside the container to see if the container is hitting its hard limit.
3. **Analyze Native Memory Tracking (NMT):** Enable `-XX:NativeMemoryTracking=summary` to identify if the memory leak is occurring in the heap or in native memory segments (e.g., excessive thread allocation).
4. **Benchmark:** Always performance-test with a memory profiler before finalizing your percentage values to ensure the off-heap headroom is sufficient for peak load.

OpenShift and Kubernetes Integration

In containerized environments, Java applications rely on integration with OpenShift and Kubernetes to manage resource allocation through cgroups. Understanding how these platforms govern resources is critical for stable deployments:

- **Resource Specifications:** Kubernetes uses `requests` and `limits` defined in deployment YAML to manage compute resources. `Requests` guide the scheduler for node placement, while `limits` enforce hard constraints on CPU and memory usage.
- **Cgroup Enforcement:** Platform resource definitions translate directly into Linux kernel cgroup values. CPU limits define execution quotas, while memory limits act as the boundary for the container's total memory footprint.
- **QoS Classes:** OpenShift employs Quality of Service (QoS) classes to handle resource contention. While "Guaranteed" QoS provides stable resource allocation, it does not prevent cgroup-level OutOfMemory (OOM) kills if the application exceeds its assigned memory limits.
- **Operational Alignment:** Proper integration requires aligning Java heap configurations with container-level constraints to prevent discrepancies between the JVM's view of memory and the actual limits imposed by the OpenShift platform.

Configuring Limits and Requests in Red Hat OpenShift

In the context of Red Hat OpenShift Container Platform (OCP), effectively managing resources is a balancing act between scheduling efficiency and runtime performance. To optimize Java applications, it is critical to distinguish how OCP handles resource definitions and how those definitions translate to the Linux kernel and the Java Virtual Machine (JVM).

The Relationship Between Requests and Limits

OpenShift utilizes Kubernetes-native resource management, defined within the container specification of a Deployment or Pod YAML.

- **Requests:** These are used by the OpenShift scheduler to place a Pod on a node. They represent the minimum amount of CPU and memory guaranteed to the container. If a node does not have enough unallocated resources to satisfy the request, the scheduler will not place the Pod on that node.
- **Limits:** These define the hard ceiling for resource consumption. If a container exceeds its memory limit, it risks being terminated by the kernel (OOMKilled). CPU limits act as a throttling mechanism, restricting the execution time available to the container within a defined period.

Resource Allocation Table

Resource Type	Purpose	Calculation Method
CPU Request	Scheduling and node assignment.	Sets <code>cpu.share</code> in cgroups to prioritize usage.

Resource Type	Purpose	Calculation Method
CPU Limit	Binding execution allocation (throttling).	Sets quota and period (defaults to 100ms segments).

Perspectives on Resource Awareness

It is vital to understand that different layers of your infrastructure interpret these settings differently:

1. **OpenShift/Kubernetes Layer:** The scheduler uses **requests** to manage cluster density and node capacity. It ignores **limits** for scheduling decisions.
2. **Kernel Layer:** The kernel enforces **limits** via Control Groups (cgroups). Quotas and periods derived from the container spec are translated into filesystem values in `/sys/fs/cgroup`.
3. **Java Layer (OpenJDK):** Modern versions of OpenJDK (post-8u191) are "container-aware." However, **Java only considers limits, not requests**. Even if you define CPU requests in your YAML, the JVM will ignore them for internal heap sizing and thread pool calculations.



While Java workloads are burstable (threads dynamically demand kernel time), the JVM performs its cgroup reading at **startup**. Consequently, the JVM is inelastic. If the user changes limits at runtime without restarting the Pod, the JVM will **not** automatically adjust its internal heuristics to match the new constraints.

Best Practices for OCP Configurations

- **Avoid Overcommitting Critical Workloads:** For production-critical applications, consider "Guaranteed" Quality of Service (QoS). A pod is considered "Guaranteed" when both requests and limits are set and equal for all containers in the pod.
- **Limit vs. Request Parity:** If you are unsure of the application's baseline, start by setting requests lower than limits to allow for burstable performance, but ensure the limit is high enough to accommodate peak load to prevent CPU throttling.
- **Memory Awareness:** Always set memory limits to prevent a single Java container from consuming all available node memory, which would trigger the node-level OOM killer and impact other pods on the same node.

Hands-on Activity: Inspecting cgroup Resource Constraints

In this activity, you will verify how OpenShift translates your YAML configuration into cgroup settings visible to the Java application.

Prerequisites: * An active OpenShift cluster. * `oc` CLI authenticated with administrative or project-level access.

Steps:

1. **Deploy a sample Java application** with defined limits and requests:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: java-resource-demo
spec:
  template:
    spec:
      containers:
      - name: java-app
        image: registry.access.redhat.com/openjdk/openjdk-11-rhel8
        resources:
          requests:
            memory: "256Mi"
            cpu: "500m"
          limits:
            memory: "512Mi"
            cpu: "1000m"
```

2. **Access the running Pod** to inspect the cgroup limits:

```
oc rsh <pod-name>
```

3. **Check the CPU quota** enforced by the kernel:

```
cat /sys/fs/cgroup/cpu/cpu.cfs_quota_us
# This will return 100000 (100ms), confirming the 1000m CPU limit.
```

4. **Observe Java's internal detection** of these limits:

```
java -XshowSettings:vm -version
```

Look for the output related to **MaxHeapSize** and **ActiveProcessorCount**. You will notice these values are derived from the limits, not the requests.

Understanding cgroups implications for Java applications

In the context of modern container orchestration, the Linux kernel's **control groups (cgroups)** feature acts as the fundamental mechanism for resource isolation. For Java applications running in OpenShift or Kubernetes, cgroups provide the enforcement boundaries for CPU and memory usage.

Understanding how the Java Virtual Machine (JVM) interacts with these cgroups is critical for preventing performance degradation and application crashes.

The Role of cgroups in Resource Management

Cgroups allow the kernel to organize processes into hierarchical groups and allocate resources—such as CPU time, system memory, and network bandwidth—to these groups. When you define **limits** and **requests** in your OpenShift Deployment configuration, OpenShift translates these values into cgroups filesystem entries (typically found under `/sys/fs/cgroup/`).

CPU and Memory Enforcement

- **CPU Limits:** Java applications are "inelastic" regarding CPU detection. Upon startup, the JVM queries the cgroup filesystem to determine how many CPU cores are available. It uses this information to size its internal thread pools (e.g., the Common ForkJoinPool). Once started, this value generally remains fixed. If a container is restricted via a quota, the kernel throttles the CPU cycles available to the JVM.
- **Memory Limits:** The JVM monitors the `memory.limit_in_bytes` (cgroups v1) or `memory.max` (cgroups v2) files. If the JVM is container-aware, it uses these values to calculate default heap sizing (e.g., typically 1/4 of the detected container memory if no explicit `-Xmx` is set).

The Danger of Decoupling

The most significant risk for Java applications in containers is the **decoupling of memory settings** from container limits. This occurs when the JVM is either non-container-aware (older versions) or is manually configured with fixed values that ignore the container's environment.

When the JVM heap size (`-Xmx`) is set higher than the actual memory limit enforced by the cgroup:

1. **The Invisible Limit:** The JVM believes it has more memory available than the container actually provides.
2. **Kernel Intervention:** When the application attempts to allocate memory beyond the cgroup limit, it does not trigger a standard `java.lang.OutOfMemoryError`. Instead, the **cgroup OOM Killer** invokes the kernel to terminate the process immediately.
3. **Observability Gaps:** Because the kernel performs the termination, these events often do not appear as standard Java application exceptions or OpenShift deployment events. This leads to silent failures, node-level instability, or a crash-looping pod that provides little diagnostic information.

Best Practices for cgroup-Safe Java

To ensure your Java applications coexist healthily with cgroup limits, adopt the following strategies:

- **Utilize Container-Aware JVMs:** Always use Java 8u191+ or modern versions (Java 11, 17, 21+), which include the JDK8146115 enhancement. These versions automatically read the cgroup filesystem to adjust memory ergonomics.
- **Use Relative Sizing:** Instead of hardcoding `-Xmx` values (e.g., `-Xmx2g`), use container-aware flags like `-XX:MaxRAMPercentage`. This allows the JVM to scale its heap dynamically based on the memory limit assigned to the container (e.g., `-XX:MaxRAMPercentage=75.0` will set the heap to 75% of the container limit).

- **Account for Off-Heap Memory:** Remember that the total memory used by a container includes the heap, the Metaspace, stack space for threads, code cache, and native buffers (Direct Byte Buffers). Ensure your `-Xmx` leaves sufficient "headroom" for these non-heap segments, or the container will exceed its cgroup limit even if the heap appears healthy.



When debugging, always check the `memory.stat` file inside the container (if permissions allow) or use tools like `jcmd` to inspect current memory usage and compare them against the deployment's defined resource limits.

Navigating Kubernetes resource management for Java

In Kubernetes and OpenShift environments, resource management is governed by the orchestration layer, which enforces boundaries using Linux Control Groups (cgroups). For Java applications, the primary challenge lies in the abstraction gap: the JVM manages its own memory (heap, code cache, metaspace) and CPU (GC threads, JIT compiler threads), while Kubernetes manages the container's total resource footprint.

Resource Requests vs. Limits

Effective resource management requires a clear understanding of the difference between requests and limits:

- **Requests:** The guaranteed amount of resources (CPU/Memory) allocated to the container. The Kubernetes scheduler uses these values to place pods on nodes.
- **Limits:** The maximum threshold allowed for the container. If a Java application exceeds its memory limit, the Linux kernel OOM (Out of Memory) Killer will terminate the process.



Setting `Limits` equal to `Requests` results in a **Guaranteed Quality of Service (QoS)** class. While this provides the highest priority, it does not prevent the OOM Killer from terminating the pod if the JVM heap plus off-heap memory usage exceeds the container limit.

Cgroups and the Java JVM

Since the introduction of JDK 8u191 (and the backport of `JDK8146115`), the JVM has become "container-aware." It reads cgroup filesystem metrics to determine available resources rather than querying the host OS kernel.

When the JVM starts, it calculates heap ergonomics based on the container's cgroup limits: * **Memory:** If no `-Xmx` is explicitly set, the JVM defaults to a fraction of the container's memory limit (typically 1/4 of the available RAM). * **CPU:** The JVM detects the number of available CPUs via cgroup shares/quotas and scales the number of Garbage Collection (GC) threads accordingly.

The Impact of Dynamic Resizing

While recent Kubernetes versions (1.27+) and OpenShift 4.4+ introduced features to resize pod resources without a restart, the JVM behavior remains tied to its initialization state.

1. **Startup Snapshot:** The JVM performs its initial heap and thread pool calculations based on the resource values present when the process is invoked.
2. **Detection vs. Adaptation:** While the JVM can technically observe changes in the cgroup filesystem during runtime, it does not necessarily resize its internal heap boundaries or thread pools dynamically.
3. **Risk:** If a container is vertically scaled (e.g., memory limit increased), the JVM will likely continue to respect the original heap size calculated at boot, potentially leaving the application under-provisioned despite the increase in available container resources.

Best Practices for Java on Kubernetes

To ensure stability in a Kubernetes-managed environment, adhere to the following configurations:

- **Explicit Heap Definition:** Even with container awareness, it is best practice to define memory settings using percentages of the container limit (e.g., `-XX:MaxRAMPercentage=75.0`). This ensures the heap scales proportionally if the container limit is adjusted in the deployment manifest.
- **Account for Off-Heap Memory:** Remember that the total container memory limit must accommodate the Java heap, Metaspace, Code Cache, Thread stacks, and Native memory (NMT). Always set the `-XX:MaxRAMPercentage` conservatively (e.g., 60-75%) to prevent OOM events.
- **CPU Throttling:** Be aware that setting strict CPU limits can lead to "CPU throttling," where the kernel pauses threads when the quota is reached. Monitor `container_cpu_cfs_throttled_periods_total` in Prometheus to identify performance bottlenecks.



Always use the latest stable version of Red Hat's OpenJDK, as it includes the most robust cgroup v2 support and the latest optimizations for container environment interactions.

Troubleshooting and Tooling

Troubleshooting Java in containerized environments requires distinguishing between tools that report host-level node metrics versus those that are container-aware. Improper resource visibility often leads to memory misconfiguration and `OOMKills`.

- **Host-level vs. Container-aware Tools:** Standard Linux utilities like `lscpu`, `nproc`, and `free -m` reflect the underlying OpenShift node's resources rather than container-specific limits.
- **Container-aware Diagnostics:** Utilize JVM-specific flags such as `VM.info` and `-XX:+PrintFlagsFinal` (via `--XshowSettings:system`) to accurately retrieve the memory and CPU constraints enforced on the container.
- **Deployment Inspection:** Analyze pod logs to verify if scripts like `run-java.sh` are active or if the default OpenJDK memory allocation is scaling correctly.
- **Heap Analysis:** Compare the JVM's reported heap size against container limit values to identify potential discrepancies between configured `-Xmx` values and actual resource availability.
- **Resource Tuning:** Be aware that manually overriding `-Xmx` on container-aware JVMs can cause mismatches with container limits, potentially leading to container underutilization or runtime `OutOfMemoryErrors`.

Identifying Container-Aware versus Non-Container-Aware Tools

When managing Java applications in Red Hat OpenShift or Kubernetes, understanding whether your diagnostic tools perceive the underlying host node or the isolated container environment is critical. Misinterpreting these outputs often leads to "out of memory" (OOM) kills or CPU throttling because the application may be configured based on node-level visibility rather than container-level constraints.

Distinguishing Tool Behavior

To effectively troubleshoot Java in containers, you must categorize your diagnostic tools based on their ability to parse `cgroups` (control groups).

Non-Container-Aware Tools

Non-container-aware tools report information regarding the physical or virtual host (the OCP Node). Relying on these tools within a pod will provide a false sense of resource availability.

- `lscpu`: Reports the total CPU architecture and count of the physical worker node, not the specific CPU quota assigned to your pod.
- `nproc`: Returns the number of processors available to the host OS, ignoring the Kubernetes `resources.limits.cpu` defined in your pod specification.
- `free -m`: Displays the total RAM available to the node. If your pod has a 512MB limit but the node has 64GB, this command will inaccurately suggest you have 64GB available for your heap.

Container-Aware Tools

Container-aware tools interface with the Linux `cgroups` subsystem to identify the specific constraints imposed by the orchestrator. These are the "source of truth" for your application environment.

- `-XX:+PrintFlagsFinal` / `-XshowSettings:system`: These flags instruct the JVM to report the memory and CPU settings it has calculated based on container limits. Use these to verify if the JVM correctly identified the container memory boundary.
- `VM.info` (via `jcmd`): Provides a snapshot of the JVM, including flags and internal settings that reflect whether the container-aware logic was successfully applied.
- `run-java.sh`: Often used in Red Hat-supported images, this script evaluates the container environment and environment variables to set optimal Java heap sizes automatically before the JVM starts.

Diagnostic Strategy

To verify if your environment is correctly configured, compare the output of host-level tools against JVM-internal tools.

Tool	Scope	Usage Case
<code>free -m</code>	Node	Use only to verify the underlying node capacity.
<code>nproc</code>	Node	Ignore when calculating thread pools or garbage collection threads.
<code>java -XshowSettings:system</code>	Container	Use to confirm the JVM sees the pod's <code>limits.memory</code> .
<code>jcmd <pid> VM.info</code>	Container	Use for deep runtime inspection of detected resource limits.

Detecting Misconfiguration

If your application is crashing with an `OOMKiller` signal, follow this diagnostic workflow:

1. **Check JVM visibility**: Run `java -XshowSettings:system` inside the running container.
2. **Compare limits**: Look at `MaxHeapSize` and `TotalPhysicalMemory` in the output.
3. **Evaluate OCP Manifests**: Compare these values against your OpenShift `Deployment` or `DeploymentConfig` resource requests and limits.
4. **Log Analysis**: Review the pod startup logs. If the logs show that the JVM is using a default (e.g., 25% of physical memory) rather than a percentage of the container limit, the JVM is likely not properly detecting the `cgroup` constraints.



If your application logs indicate it is defaulting to 25% of the **node's** memory rather than the **container's** limit, verify that you are using a supported OpenJDK

version (8u191+) and that the image is properly configured to handle **cgroups**.

Analyzing Container Limits and Resource Utilization

To ensure Java applications perform predictably within Red Hat OpenShift or Kubernetes, you must understand how the JVM interacts with the underlying control groups (cgroups). When an application is container-aware, the JVM can query the cgroup filesystem to determine its actual boundaries rather than defaulting to the physical host's resources.

Understanding Resource Visibility

The primary challenge in containerized Java is the discrepancy between what the Operating System sees (physical host) and what the container runtime enforces (cgroup limits).

The Mechanics of Detection

Container-aware JVMs (starting from Java 8u191 and newer) automatically inspect cgroup files—such as **memory.limit_in_bytes** and **cpu.cfs_quota_us**—to size the heap and determine the available processor count.

When analyzing your application, you can extract this metadata using **jcmd**. If your JVM is properly detecting the container environment, the output will reveal specific cgroup configurations:

```
# Example output from jcmd <pid> VM.check_commercial_features or custom diagnostic
commands
container_type: cgroupv1
active_processor_count: 1
cpu_quota: -1
memory_limit_in_bytes: 536870912
memory_usage_in_bytes: 125108224
```

- **active_processor_count**: Essential for the JVM to size its internal thread pools (e.g., ForkJoinPool).
- **memory_limit_in_bytes**: Dictates the physical memory limit the JVM uses to calculate its initial heap size (typically 1/4 of the container limit by default).

Identifying Limitations and Misconfigurations

Analyzing resource utilization requires verifying that your JVM heap plus off-heap requirements remain within the Kubernetes/OpenShift **limits**.

The Decoupling Risk

A common pitfall is ignoring the "off-heap" footprint. If you set your Java Heap (Xmx) to 80% of your container limit, but your application requires 30% for off-heap memory (Metaspace, Code Cache, Thread stacks, Direct Buffers), the container will exceed its memory limit.

This results in a **cgroup OOMKill**, which is distinct from a JVM `java.lang.OutOfMemoryError`.

Troubleshooting Checklist

When analyzing resource constraints, follow these steps:

1. **Compare Requests vs. Limits:** Ensure your OpenShift `requests` are sufficient for startup, while `limits` accommodate peak load.
2. **Verify Cgroup Awareness:** Use `jcmd` to verify if the JVM identifies the correct `memory_limit_in_bytes`. If this shows `-1` or reflects host memory, the JVM is not container-aware, and you must explicitly configure `-XX:MaxRAMPercentage`.
3. **Monitor via Tools:**
 - **Container-aware:** Use `kubectl top pod` or the OpenShift Web Console metrics.
 - **Non-container-aware:** Avoid using legacy `top` or `free` commands inside the container, as they often report host-level data rather than container-constrained data.

Hands-on Activity: Inspecting JVM Resource Awareness

In this activity, you will verify if your currently running pod is correctly interpreting its cgroup limits.

1. **Access the running pod:**

```
oc rsh <pod-name>
```

2. **Identify the PID of the Java process:**

```
jps
# Output: 1234 AppMain
```

3. **Run the diagnostic command to check resource visibility:**

```
jcmd 1234 VM.native_memory summary
# Look for container-specific configuration in the JVM diagnostic output
```

4. **Cross-reference with OpenShift limits:** Compare the `memory_limit_in_bytes` reported by the JVM with the memory limit defined in your OpenShift deployment YAML. If they do not align, adjust your `JAVA_TOOL_OPTIONS` using `-XX:MaxRAMPercentage` to force the JVM to align its heap to the container boundary.



If you observe **cgroup OOMKills** despite the JVM reporting correct memory usage, analyze your off-heap usage via `jcmd <pid> NativeMemoryTracking=summary`. You may need to decrease your `MaxRAMPercentage` to provide more "breathing room" for the native memory footprint.

Diagnostic strategies for Java deployments in OpenShift

When managing Java applications in OpenShift, traditional diagnostic tools often fail because they lack "container awareness." Tools that look at the host operating system's total resources rather than the container's constrained cgroup limits can lead to misleading metrics. Effective diagnostics require tools and strategies that respect the isolation boundaries defined by Kubernetes.

Identifying Container-Aware Tools

A critical diagnostic step is distinguishing between container-aware and legacy tools.

- **Legacy Tools:** Tools like `top`, `free`, or older JVM monitoring utilities may report the host's total RAM and CPU count. Relying on these leads to "false positives" where an administrator believes the application has more resources than it actually does.
- **Container-Aware Tools:** Modern JDKs (8u191+, 11+) and utilities like `jcmd`, `jstat`, and modern observability agents are cgroup-aware. They query the container's `/sys/fs/cgroup/` filesystem to report limits specifically applied to the pod.

Diagnostic Strategies

To effectively diagnose resource constraints or performance issues, follow these strategies:

1. Validating Resource Visibility

Always verify that the JVM sees the limits defined in your OpenShift `DeploymentConfig` or `Deployment`. You can inspect this by executing into the pod and checking the container's perceived memory and CPU:

```
# Check perceived max heap from within the pod
oc exec <pod-name> -- java -XshowSettings:vm -version
```

Check if `MaxHeapSize` is calculated correctly based on the container memory limit rather than the worker node's total memory.

2. Utilizing `jcmd`

`jcmd` is the gold standard for Java diagnostics within OpenShift. It allows you to perform heap dumps, thread dumps, and analyze VM flags without requiring a restart.

- **Thread Dumps:** Identify if the application is hitting thread limits or experiencing deadlocks.
- **GC Performance:** Use `jcmd <pid> GC.run` to trigger garbage collection or `jcmd <pid> VM.flags` to verify if `-XX:MaxRAMPercentage` is correctly configured.

3. Analyzing cgroup Constraints

If a pod is frequently restarting, it may be hitting OOM (Out of Memory) kills. Since the OOM Killer

is enforced at the cgroup level, standard Java `OutOfMemoryError` logs might not appear if the kernel terminates the process before the JVM can report it.

- **Inspect Pod Events:** Use `oc describe pod <pod-name>` to check for the `OOMKilled` state.
- **Check Exit Codes:** An exit code of `137` is a strong indicator that the container was terminated by the OOM killer because it exceeded the OpenShift memory limit.

Hands-on Lab: Analyzing Resource Constraints

In this lab, we will use `jcmd` to inspect a running Java process and verify its awareness of container limits.

Prerequisites

- An active OpenShift cluster with a running Java pod.
- `oc` CLI configured with cluster access.

Step 1: Verify JVM resource settings

Execute into the pod to check how the JVM views the memory limits.

```
# Run the VM settings check
oc exec <pod-name> -- java -XX:+PrintFlagsFinal -version | grep MaxHeapSize
```

Observe the output. If you have defined a `memory` limit of 512Mi in your manifest, the `MaxHeapSize` should be a fraction of that limit (if using `MaxRAMPercentage`).

Step 2: Use `jcmd` to analyze the process

Access the container shell to utilize `jcmd`.

```
oc rsh <pod-name>

# List running Java processes
jcmd

# Get a detailed report of the VM's perceived memory constraints
jcmd <pid> VM.flags
```

Step 3: Troubleshooting a constraint

If the application is slow, check if CPU throttling is occurring due to the Kubernetes `cpu` limit:

1. Check the `cpu.stat` file within the container: `[source,bash] ---- cat /sys/fs/cgroup/cpu/cpu.stat ----`
2. Look for `nr_throttled`. If this number is increasing rapidly, your CPU `limit` is too low, and the kernel is throttling your application, causing latency.

Best Practices Summary

- **Never rely on `free` or `top`** inside a container for memory capacity planning.
- **Always use `-XX:MaxRAMPercentage`** instead of fixed `-Xmx` values to ensure the heap scales dynamically with the container limit.
- **Prioritize `jcmd`** for all runtime analysis, as it is native to the OpenJDK and fully supports container-cgroup awareness.

Java Diagnostics: `jcmd`

The `jcmd` utility is a versatile diagnostic tool included with the JDK. In containerized environments, it is indispensable for inspecting the internal state of a running Java Virtual Machine (JVM) without requiring external profilers or complex agent attachments.

Role of Java Diagnostics (`jcmd`) in Containerized Environments

In a containerized world, `jcmd` allows developers and operators to bridge the gap between the container's resource limits (cgroups) and the JVM's internal resource allocation. When troubleshooting performance or memory issues in OpenShift or Kubernetes, `jcmd` provides a direct view into the JVM's view of the container.

Key Capabilities

- **Process Identification:** List running Java processes within the container.
- **Performance Diagnostics:** Retrieve thread dumps, GC statistics, and class loading metadata.
- **Memory Management Analysis:** Inspect heap usage and native memory tracking (NMT).
- **Dynamic Configuration:** Adjust certain JVM flags at runtime (where supported).

Using `jcmd` for Resource Inspection

The main tool for Java investigations inside a container is `VM.info`, which is available in JDK 11+. It prints several details on the Java process itself, from the JVM perspective. See full review on [Interpreting VM.info file in OpenJDK/OracleJDK](#).

When a container is suffering from `OutOfMemory` (OOM) errors or unexpected latency, you can execute `jcmd` directly inside the pod to gather data.

```
# List all running Java processes
jcmd

# Java full dump (JDK 11+)
jcmd <PID> VM.info

# Retrieve a thread dump for a specific PID
jcmd <PID> Thread.print

# Check GC performance statistics
```

```
jcmd <PID> GC.run
```

Analyzing Native Memory

One of the most critical aspects of containerized Java is distinguishing between the **Java Heap** and **Off-Heap memory**. If the container is being killed by the OOMKiller but the heap appears healthy, the issue often resides in native memory (e.g., Metaspace, code cache, or direct buffers).

Enable Native Memory Tracking (NMT) via the JVM argument `-XX:NativeMemoryTracking=summary` to use `jcmd` to identify leaks:

```
# View native memory summary
jcmd <PID> VM.native_memory summary
```

Best Practices for Diagnostics in OpenShift

1. **Accessing the Pod:** Use `oc rsh <pod-name>` to gain shell access to the running container.
2. **Tooling Availability:** Ensure your container base image includes the `jcmd` binary (standard in most OpenJDK-based images).
3. **Non-intrusive Data Collection:** Use `jcmd` for point-in-time snapshots. Avoid high-frequency polling, as diagnostics can consume CPU cycles within the container's limited quota.
4. **Verification:** Compare the output of `GC.heap_info` against your pod's `resources.limits.memory` to verify if the JVM is respecting the container's memory boundaries.

Hands-on Lab: Inspecting JVM Resources

In this exercise, we will verify the JVM's memory awareness.

1. **Access the container terminal:**

```
oc rsh deployment/my-java-app
```

2. **Identify the PID of the Java process:**

```
jcmd -l
```

3. **Inspect the current heap configuration:** Use the following command to check how the JVM is interpreting the container memory limits:

```
jcmd <PID> VM.flags
```

Look for `-XX:MaxHeapSize` or `-XX:InitialHeapSize`.

4. **Analyze GC performance:** Check if the Garbage Collector is struggling due to container

constraints:

```
jcmd <PID> GC.class_histogram
```

Troubleshooting Tip: If `jcmd` reports a heap size significantly different from your OpenShift memory limits, check your deployment YAML to ensure you are not manually overriding the heap size with `-Xmx`, which might be conflicting with the container-aware auto-tuning.

Hands-on Lab

The hands-on lab focuses on practical application of container awareness principles within an OpenShift environment. Key activities include:

- **Resource Tuning:** Adjusting Java heap settings using percentage-based configurations to align with container-imposed limits.
- **Monitoring:** Utilizing container-aware diagnostic tools to track actual CPU and memory utilization rather than host-level metrics.
- **Troubleshooting:** Applying diagnostic strategies—such as inspecting pod configurations and resource constraints—to resolve deployment issues within OpenShift.

Tuning Java Heap Settings Based on Container Limits

In modern containerized environments like Red Hat OpenShift, the JVM no longer treats the host machine's total physical memory as its boundary. Instead, it respects the **cgroup** limits assigned to the container. Tuning the Java heap using percentages is the industry-standard approach to ensure the JVM remains "elastic" relative to the container's resource allocation.

The Shift to Percentage-Based Memory Management

Historically, developers used the **-Xmx** flag to set an absolute heap size (e.g., **-Xmx2g**). In a container, this creates a static configuration that ignores the actual memory limit of the pod. If the container limit is lowered, the JVM may attempt to exceed it, leading to a **cgroup** OOM-Kill.

By using **RAM Percentage flags**, you decouple the specific memory limit from the application code, allowing the JVM to calculate its heap size dynamically based on the container's available memory.

Key Configuration Flags

- **-XX:MaxRAMPercentage=N**: Sets the maximum heap size as a percentage of the container's memory limit. The default is typically 25%, though specialized container images (like Red Hat's **run-java.sh** scripts) often default to 50% or 80%.
- **-XX:InitialRAMPercentage=N**: Sets the initial heap size as a percentage of the container limit.
- **-XX:MinRAMPercentage=N**: Defines the minimum heap size used when the container memory limit is small.



When tuning these percentages, remember that the **total container memory** must accommodate more than just the heap. You must account for **off-heap memory**, which includes:

- Metaspace (Class metadata)
- Thread stacks
- Code Cache
- GC structures and internal JVM native memory

- Off-heap buffers (used by Netty, AMQ, or Data Grid)

Setting `-XX:MaxRAMPercentage` too high (e.g., 95%) leaves insufficient room for native memory, inevitably leading to native memory pressure and `cgroups` OOM-Kill.

Hands-on Lab: Tuning Heap for a Containerized Application

In this lab, you will configure a Java application to use container-aware memory settings instead of hardcoded heap limits.

Prerequisites

- Access to an OpenShift cluster or a local Docker environment.
- A containerized Java application image.

Step 1: Inspect Current Configuration

First, verify that your application is currently running without explicit `-Xmx` flags or is using percentage-based flags. Check the environment variables or the deployment configuration:

```
# Check the deployment environment
oc set env deployment/<deployment-name> --list
```

Step 2: Configure Percentage-Based Heap

Update your deployment configuration to utilize `MaxRAMPercentage`. This ensures that if you increase the pod's memory limit in the future, the heap scales automatically without requiring a code or environment variable change.

Apply the following environment variables to your deployment:

```
oc set env deployment/<deployment-name> \
  JAVA_OPTIONS="-XX:MaxRAMPercentage=75.0 -XX:InitialRAMPercentage=50.0"
```

Step 3: Validate the Configuration

After the deployment restarts, verify that the JVM has correctly calculated the heap size relative to your pod's memory limit.

1. Get the pod name:

```
oc get pods
```

2. Run `jcmd` inside the container to inspect the VM arguments:

```
oc exec <pod-name> -- jcmd 1 VM.flags
```

Observe the output: You should see the `-XX:MaxRAMPercentage` flag active. If you check `VM.max_heap_size`, you will see an absolute byte value calculated based on the 75% ratio of your current container limit.

Step 4: Monitoring

Use the OpenShift web console or CLI to monitor memory usage. Compare the "Memory Usage" metric against the "Container Limit." Ensure that the "Usage" remains below the total limit, accounting for the gap required for native off-heap memory.



If you observe `java.lang.OutOfMemoryError: Metaspace` or native crashes, your `MaxRAMPercentage` is too aggressive. Lower the value by 5-10% and redeploy.

Monitoring CPU and Memory Usage with Container-Aware Tools

When deploying Java applications in containerized environments like Red Hat OpenShift or Kubernetes, traditional Linux diagnostic tools often fall short. Tools that are not "container-aware" report the physical or virtual resources of the entire host node rather than the specific constraints applied to your pod.

This section explores how to distinguish between container-aware and node-aware outputs and how to use them effectively for troubleshooting.

The Visibility Gap

If you use standard Linux utilities such as `lscpu`, `nproc`, or `free -m` inside a container, you are likely seeing the resources of the underlying host node. Relying on these values for capacity planning can lead to "Out of Memory" (OOM) kills or CPU throttling because the JVM may attempt to allocate resources exceeding the container's cgroup limits.

To monitor your Java application accurately, you must utilize tools and flags that interface directly with the `cgroups` filesystem.

Container-Aware Diagnostic Tools

1. JVM Flags: `-XX:+UnlockDiagnosticVMOptions`

The JVM provides built-in mechanisms to report exactly what it sees regarding container limits.

- `--XshowSettings:system`: This is the most critical flag for initial diagnosis. It displays the system information the JVM has detected, including cgroup-enforced memory limits and the active processor count.
- `VM.info`: When using `jcmd`, this command provides a comprehensive dump of the current JVM

configuration, including container-specific constraints.

2. Interpreting Container-Aware Output

When running a container-aware command (such as `java -XshowSettings:system -version`), look for the following key indicators:

- **active_processor_count**: Represents the number of CPUs available to the JVM based on the Kubernetes/OpenShift CPU limits.
- **memory_limit_in_bytes**: Reflects the specific cgroup memory limit. If this value does not match your expected Kubernetes resource limit, the JVM is not correctly detecting the container boundary.
- **cpu_quota** and **cpu_period**: These values dictate the CPU throttling behavior. A **cpu_quota** of `-1` suggests no limit is set, while specific values indicate restricted processing time.

Hands-on Activity: Verifying Container Limits

In this activity, you will compare "Node-aware" vs "Container-aware" data to ensure your JVM is correctly bounded.



Ensure you have a running Java pod in your OpenShift/Kubernetes environment and access to `oc rsh` or `kubectl exec`.

Step 1: Compare Node vs. Container Visibility

1. Log into your running Java pod:

```
oc rsh <pod-name>
```

2. Run a non-container-aware command to see the node's resources:

```
free -m  
nproc
```

Observe: The output reflects the total RAM and CPUs of the entire OpenShift worker node.

3. Run the container-aware JVM setting command:

```
java -XX:+UnlockDiagnosticVMOptions -XshowSettings:system -version
```

Step 2: Analyze the JVM Output

Locate the container information section in the output. You should see entries similar to:

```
container (cgroup) information:
```



```
container_type: cgroupv1
cpu_cpuset_cpus: 0-1
active_processor_count: 2
memory_limit_in_bytes: 536870912
```

- **Verification:** Confirm that `memory_limit_in_bytes` aligns with the memory `limit` defined in your pod's YAML configuration.
- **Action:** If these numbers differ from your configured limits, check if your JVM version is 8u191 or later, as older versions lack proper cgroup awareness.

Step 3: Real-time Analysis with jcmd

1. While the application is running, use `jcmd` to inspect the live process:

```
jcmd <pid> VM.info
```

2. Search for the `cgroup` section in the output. This provides a real-time snapshot of the container metrics, including `memory_usage_in_bytes`. Use this to determine if your heap size (typically a percentage of the limit) is appropriately leaving enough headroom for off-heap memory.

Troubleshooting Resource Constraints in OpenShift

When Java applications run in OpenShift, they operate within the constraints defined by the Linux kernel's control groups (cgroups). If these constraints are misconfigured or if the Java Virtual Machine (JVM) is unaware of them, you may encounter performance degradation, unexpected `OutOfMemoryError` (OOM) exceptions, or even container restarts due to OOMKills.

This section focuses on diagnostic strategies and using tools to inspect resource constraints for Java workloads in OpenShift.

Diagnostic Strategy

To effectively troubleshoot resource constraints, you must correlate three distinct layers:

1. **OpenShift/Kubernetes Layer:** The resource requests and limits defined in your `Deployment` or `Pod` specification.
2. **Kernel/Cgroup Layer:** The actual hardware quotas enforced by the container runtime.
3. **JVM Layer:** The memory and CPU settings perceived by the Java process.

Identifying Misalignment

A common pitfall is the misalignment between the container limit and the JVM heap size. If your container limit is 1Gi but the JVM is configured with `-Xmx2G`, the process will likely fail immediately or fail shortly after starting, being killed by the kernel when it attempts to allocate memory beyond the cgroup boundary.

Tools for Inspection

Using `oc describe`

The first step in any troubleshooting process is verifying the resource configuration of the pod.

```
oc describe pod <pod-name>
```

- **Look for:** The `Limits` and `Requests` sections.
- **Action:** Ensure the `Limits` are explicitly defined. If `Limits` are missing, the JVM may default to host-level visibility rather than container-level, potentially leading to incorrect heap sizing.

Leveraging `oc debug` and `inspect`

When an application behaves unexpectedly, you can use `oc debug` to inspect the runtime environment from inside the pod's namespace.

```
# Access the pod environment
oc debug pod/<pod-name> --image=registry.access.redhat.com/ubi9/ubi-minimal

# Once inside, inspect the cgroup limits directly
cat /sys/fs/cgroup/memory/memory.limit_in_bytes
```

Note: In cgroup v2 environments, the path may be located under `/sys/fs/cgroup/memory.max`.

Using `jcmd`

The `jcmd` utility is the standard diagnostic tool for Java. Since Java 8u191+, the JVM is container-aware. You can use `jcmd` to verify how the JVM "sees" the container.

```
# Run jcmd inside the container to see VM flags
oc exec <pod-name> -- jcmd 1 VM.flags
```

- **Look for:** `MaxHeapSize`
- **Verification:** If `MaxHeapSize` is disproportionately large compared to your container limit, your application is at risk of being killed by the kernel.

Hands-on Lab: Troubleshooting a Resource-Constrained Pod

In this exercise, you will diagnose a pod experiencing memory constraints.

Prerequisites:

- An active OpenShift cluster.
- The `oc` CLI tool authenticated to your cluster.

Steps:

1. **Simulate a kill:** Deploy a pod with a very low memory limit (e.g., 64Mi) and a JVM heap requirement that exceeds directly set via Xmx. Observe the kill:

```
[source,bash]
----
oc get pods
# Look for OOMKilled status
----
```

2. **Inspect the exit code:**

```
[source,bash]
----
oc get pod <pod-name> -o
jsonpath='{.status.containerStatuses[0].state.terminated.reason}'
----
*If the result is `OOMKilled`, the kernel terminated the process because it
exceeded the cgroup memory limit.*
```

3. **Verify the limits:** Check the deployment configuration to see the discrepancy between your **limit** and the JVM's **-Xmx** setting.
4. **Remediate:** Adjust the deployment YAML to either increase the container memory limit or update the JVM flags using the **JAVA_OPTS** environment variable to reflect a lower heap percentage:

```
[source,bash]
----
env:
  - name: JAVA_OPTS
    value: "-XX:MaxRAMPercentage=50.0"
----
```

5. **Re-verify:** After applying the change, use **jcmd** inside the pod to confirm the new **MaxHeapSize** is calculated correctly based on the updated container limit.

Appendix

FIXME: Content for Appendix...

Resources

[Download Course PDF](#)